

# COMPUTER NETWORK

## ASSIGNMENT 2

### REPORT

#### QUESTION 1

##### Server Code

- 1) The server sets up a TCP socket and listens for client connections.

##### Socket Creation:

- The **socket()** function creates a TCP socket (with AF\_INET for IPv4 and SOCK\_STREAM for TCP).
- If the socket creation fails, an error message is printed using **perror()**, and the program exits.

##### Binding:

- The **bind()** function associates the server socket with a specific IP address and port, allowing it to listen on that address.
- If the binding fails, the program prints an error message and exits.

##### Listening:

- The **listen()** function puts the server in a passive listening mode, allowing it to accept incoming client connections. The second argument specifies the maximum number of pending connections in the queue.
- If the server fails to listen, it prints an error message and exits.

- 2) The server handles multiple concurrent clients (multithreaded, concurrent server). The server accepts the client connection; hence, a new socket is created with **4 tuples** (server IP, server listening port, client IP, client port). The server creates a new thread that continues to process the client connection

##### Accepting Client Connections:

- The **accept()** function waits for incoming client connections on the server's listening socket. When a client connects, a new socket (new\_socket) is created specifically for that client connection.
- The 4-tuple created consists of:
  - Server IP
  - Server port
  - Client IP (retrieved from the client connection)
  - Client port (assigned by the client system)

- If the `accept()` call fails, an error message is printed, and the server continues to listen for the next client.

#### Thread Creation:

- For each new client, the server creates a new thread using **`pthread_create()`**. This allows the server to handle multiple clients concurrently.
- The `new_socket` is passed to the thread by allocating memory for the socket pointer and passing it as an argument to the **`handle_client()`** function.

#### Thread Detachment:

- **`pthread_detach()`** is called to ensure the thread cleans up its resources when it finishes. This allows the server to handle multiple client threads without manual thread management.
- 3) The original server socket continues to listen on the same listening port for newer incoming client connections.
    - The server is placed in an infinite loop, allowing it to call `accept()` and handle new client connections continuously.
    - Each time `accept()` is called, the server waits for a new client connection on the same listening port (8080). Once a client connects, the server creates a new socket (`new_socket`) for that specific connection, but the original socket (`server_fd`) continues listening for additional clients.

#### Non-Blocking Server:

- After each client connection is accepted and handed off to a new thread via **`pthread_create()`**, the server immediately returns to the top of the loop and waits for the next client connection.
  - This allows the original server socket (`server_fd`) to remain open on the same port and handle new incoming connections while other clients are being processed in separate threads.
- 4) The server finds out the top CPU-consuming process (user+kernel CPU time) and gathers information such as process name, pid (process id), and CPU usage of the process in user & kernel mode
    - This functionality is handled by the **`get_process_info()`** function, which reads the `/proc` directory to find running processes and gathers information such as process name, PID, and CPU usage in user and kernel modes. It calculates the total CPU time (user time + kernel time) for each process and identifies the top two CPU-consuming processes.

## CLIENT

- 1) The client creates a socket and initiates the TCP connection. The client process supports initiating “n” concurrent client connection requests, where “n” is passed as a program argument.
  - The client creates a socket using **socket()** to establish a **TCP connection**. If the creation fails, it prints an error message and exits the thread.
  - The client supports initiating “n” **concurrent connection requests** by creating multiple threads. The number of client threads is determined by the command-line argument passed when executing the program. Each thread runs the **client\_thread()** function, which handles the connection and communication with the server.
- 2) After the client connection is established, the client sends a request to the server to get information about the server’s top TWO CPU-consuming processes.
  - After the client establishes a TCP connection, it sends a message to the server, indicating a request for information about the top two CPU-consuming processes.
- 3) The client prints this information & closes the connection.
  - The **read()** function is used to receive the server's response, which is then printed to the console.
  - After the response is printed, the client closes the connection using **close()**, indicating that the communication with the server is complete.

NOTE: taskset was used for client and server code to pin task to specific CPUs

## QUESTION 2

NOTE: Readings were taken after using taskset to pin tasks to specific CPUs.

### (a) Single-threaded TCP client-server

```
ubuntu@ubuntu-dual-machine:~/Downloads/A2/q2$ sudo taskset -c 0 perf stat ./server
Server is now listening for incoming connections on port 8002
Message received from client: Requesting top CPU processes
Message received from client: Requesting top CPU processes
Message received from client: Requesting top CPU processes
Client disconnected.
^X^C./server: Interrupt

Performance counter stats for './server':

   15.39 msec task-clock                #    0.002 CPUs utilized
         7 context-switches            #   454.911 /sec
          0 cpu-migrations              #    0.000 /sec
        143 page-faults                #    9.293 K/sec
<not counted> cpu_atom/cycles/         #
26,622,471 cpu_core/cycles/            #    1.730 GHz
<not counted> cpu_atom/instructions/    #
52,089,912 cpu_core/instructions/      #
<not counted> cpu_atom/branches/       #
9,614,604  cpu_core/branches/          #   624.826 M/sec
<not counted> cpu_atom/branch-misses/  #
68,099    cpu_core/branch-misses/     #
TopdownL1 (cpu_core)                  #   30.7 % tna_backend_bound
                                           #    6.4 % tna_bad_speculation
                                           #   29.4 % tna_frontend_bound
                                           #   33.5 % tna_retiring

 6.487866455 seconds time elapsed

 0.003354000 seconds user
 0.012300000 seconds sys

ubuntu@ubuntu-dual-machine:~/Downloads/A2/q2$
```

### Performance counter stats for './server'

15.39 msec	task-clock	#	0.002 CPUs utilized
7	context-switches	#	454.911 /sec
0	cpu-migrations	#	0.000 /sec
143	page-faults	#	9.293 K/sec
26,622,471	cpu_core/cycles/	#	1.730 GHz
52,089,912	cpu_core/instructions/		
9,614,604	cpu_core/branches/	#	624.826 M/sec
68,099	cpu_core/branch-misses/		
6.487866455 seconds	time elapsed		
0.003354000 seconds	user		
0.012300000 seconds	sys		

### (b) Multi-threaded TCP client-server

```
ubuntu@ubuntu-dual-machine: ~/Downloads/A2/q2$ sudo taskset -c 0 perf stat ./server
Message received from client 0: Hello from client
Message received from client 1: Hello from client
Message received from client 2: Hello from client
^C./server: Interrupt
Performance counter stats for './server':
    22.56 msec task-clock      #    0.004 CPUs utilized
         17 context-switches  # 753.653 /sec
          0 cpu-migrations     #    0.000 /sec
        1,746 page-faults     # 77.405 K/sec
<not counted> cpu_aton/cycles/          #    3.042 GHz          (0.00%)
68,613,095 cpu_core/cycles/          #    3.042 GHz          (0.00%)
<not counted> cpu_aton/instructions/      #
106,328,589 cpu_core/instructions/      #
<not counted> cpu_aton/branches/          #    875.197 M/sec      (0.00%)
19,741,646 cpu_core/branches/          #    875.197 M/sec      (0.00%)
<not counted> cpu_aton/branch-misses/     #
233,960 cpu_core/branch-misses/     #
TopdownL1 (cpu_core)        # 34.4 % tma_backend_bound
                                #  8.5 % tma_bad_speculation
                                # 29.6 % tma_frontend_bound
                                # 27.5 % tma_retiring

6.134550201 seconds time elapsed

0.003746000 seconds user
0.008343000 seconds sys

ubuntu@ubuntu-dual-machine: ~/Downloads/A2/q2$
```

### Performance counter stats for './server'

22.56 msec	task-clock	#	0.004 CPUs utilized
17	context-switches	#	753.653 /sec
0	cpu-migrations	#	0.000 /sec
1,746	page-faults	#	77.405 K/sec
68,613,095	cpu_core/cycles/	#	3.042 GHz
106,328,589	cpu_core/instructions/		
19,741,646	cpu_core/branches/	#	875.197 M/sec
233,960	cpu_core/branch-misses/		

6.13455021 seconds

time elapsed

0.003746000 seconds

user

0.008343000 seconds

sys

(c) TCP client-server using “select”

```
ubuntu@ubuntu-dual-machine: ~/Downloads/A2/q2$ sudo taskset -c 0 perf stat ./server
Server is now listening for incoming connections.
Accepted new client connection.
Message received from client: Requesting top CPU processes
Message received from client: Requesting top CPU processes
Message received from client: Requesting top CPU processes
Client disconnected.
^C./server: Interrupt

Performance counter stats for './server':

      28.92 msec task-clock                #    0.007 CPUs utilized
           15 context-switches           #   518.673 /sec
              0 cpu-migrations            #    0.000 /sec
          1,686 page-faults               #   58.299 K/sec
<not counted> cpu_aton/cycles/            #
67,269,978   cpu_core/cycles/            #    2.326 GHz
<not counted> cpu_aton/instructions/       #
100,160,377  cpu_core/instructions/       #
<not counted> cpu_aton/branches/          #
18,528,735   cpu_core/branches/          #   640.691 M/sec
<not counted> cpu_aton/branch-misses/     #
231,798      cpu_core/branch-misses/     #
      TopdownL1 (cpu_core)               #   35.1 % tna_backend_bound
                                           #    8.6 % tna_bad_speculation
                                           #   29.9 % tna_frontend_bound
                                           #   26.4 % tna_retiring

      4.050135388 seconds time elapsed

      0.002856000 seconds user
      0.012377000 seconds sys
```

Performance counter stats for './server'

28.92 msec

task-clock

#

0.007 CPUs utilized

15

context-switches

#

518.673 /sec

0

cpu-migrations

#

0.000 /sec

1,686

page-faults

#

58.299 K/sec

67,269,978

cpu\_core/cycles/

#

2.326 GHz

100,160,377

cpu\_core/instructions/

#

18,528,735

cpu\_core/branches/

#

640.691 M/sec

231,798

cpu\_core/branch-misses/

#

4.050135388 seconds

time elapsed

0.002856000 seconds

user

0.012377000 seconds

sys

1. Task Clock and CPUs Utilized:

- **Single-Threaded Server:**
  - Task clock: 15.39 msec

- CPUs utilized: 0.002
- **Multi-threaded Server:**
  - Task clock: 22.56 msec
  - CPUs utilized: 0.004
- **Server using Select:**
  - Task clock: 28.92 msec
  - CPUs utilized: 0.007

The **Server using Select** spent the longest amount of time actively executing instructions on the CPU, while the **Single-Threaded Server** spent the least time. The **Multi-threaded Server** shows moderate execution time, which aligns with its more complex structure compared to the single-threaded version.

## 2. Context Switches (switches between kernel and user space):

- **Single Thread:** 7 context switches
- **MultiThread:** 17 context switches
- **Select-based:** 15 context switches
- The **single-threaded** server has the fewest context switches, as expected. It handles requests sequentially in a single process, leading to minimal task switching.
- The **multi-threaded** server has the most context switches because each thread may need to switch contexts frequently to handle multiple requests concurrently.
- The **select-based** server has context switches slightly lower than the multi-threaded one. The select-based model, though handling multiple connections, typically does so within a single thread, reducing the need for frequent switching but still causing more than a single-threaded model.

## 3. Page Faults:

- **Single-Threaded Server:** 143 (9.293 K/sec)
  - **Multi-threaded Server:** 1,746 (77.405 K/sec)
  - **Server using Select:** 1,686 (58.299 K/sec)
- 
- The **single-threaded** server has the fewest page faults, probably because it doesn't handle many requests simultaneously and thus has fewer memory access conflicts.
  - Both **multi-threaded** and **select-based** servers have significantly higher page faults. This makes sense, as both handle multiple requests simultaneously, potentially needing more memory allocation, leading to more page faults.

#### 4. CPU Core Cycles:

- **Single-Threaded Server:** 26,622,471 cycles (1.730 GHz)
- **Multi-threaded Server:** 68,613,095 cycles (3.042 GHz)
- **Server using Select:** 67,269,978 cycles (2.326 GHz)

The **multi-threaded** server uses the most CPU cycles, closely followed by the select-based server. This can be explained by both servers having more tasks to manage in parallel (multiple threads or connections). The **single-threaded** server, with fewer tasks at a time, uses fewer cycles overall.

#### 5. Instructions:

- **Single-Threaded Server:** 52,089,912 instructions
- **Multi-threaded Server:** 106,328,509 instructions
- **Server using Select:** 100,160,377 instructions

The **Multi-threaded Server** executes more than double the number of instructions compared to the **Single-Threaded Server**, as expected due to the complexity of handling multiple threads. The **Server using Select** executes a similar number of instructions to the multi-threaded server.

#### 6. Branch Misses:

- **Single-Threaded Server:** 68,099 branch misses
- **Multi-threaded Server:** 233,960 branch misses
- **Server using Select:** 231,798 branch misses

**Branch misses** represent inefficiency in CPU operations. The multi-threaded and select-based servers both have significantly more branch misses, which is expected because they handle more complex and parallelized logic, increasing the chance of mispredictions. The single-threaded server, with simpler flow control, has far fewer branch misses

#### 7. User and Sys Time:

- **Single-Threaded Server:**
  - User: 0.003354 seconds
  - Sys: 0.012300 seconds
- **Multi-threaded Server:**
  - User: 0.003746 seconds
  - Sys: 0.008343 seconds
- **Server using Select:**
  - User: 0.002850 seconds
  - Sys: 0.012377 seconds

#### 8. CPU Migrations: Zero for all the three as taskset pinned the task to specific CPUs

The **Single-Threaded Server** and **Select Server** spend a slightly higher amount of time in system calls (sys time), while the **Multi-threaded Server** has more time in user space (likely due to managing multiple threads).

Summary:

- **Single Threaded**: Best for simple, low-concurrency applications where system resources (e.g., memory and CPU cycles) must be minimized.
- **Multi-Threaded**: Suitable for high-load applications requiring parallel processing, though it incurs a higher system overhead.
- **Select-based**: Best for moderately loaded servers where efficient handling of multiple connections is needed without the heavy overhead of multi-threading.